

MongoDB Interview Q&As (Part 2)

◆ Section 7: Advanced Aggregation

Q31. What is \$unwind in MongoDB?

Answer: \$unwind deconstructs an array field → creates one document per array element.

```
db.orders.aggregate([{$unwind: "$items"}])
```

💡 **Tip:** Combine with \$group to analyze array data (e.g., item sales).

Q32. What is \$facet used for?

Answer: Runs multiple pipelines in a single aggregation. Example → get paginated results + counts at once.

💡 **Tip:** Useful for dashboards and reports.

Q33. What is \$project in aggregation?

Answer: Reshapes documents, includes/excludes fields, creates computed values.

```
db.users.aggregate([{$project: { name: 1, yearOfBirth: { $year: "$dob" } } }])
```

💡 **Tip:** Think of \$project = SQL's SELECT column AS alias.

Q34. What is \$match vs \$group?

Answer:

- \$match → filters documents (like WHERE).
- \$group → groups documents (like GROUP BY).

💡 **Tip:** Always use \$match before \$group for performance.

Q35. What is a pipeline in MongoDB?

Answer: A series of aggregation stages where output of one stage becomes input of the next.

💡 **Tip:** Analogy = SQL query execution plan.

◆ Section 8: Schema Design & Modeling

Q36. What is the difference between Embedded and Referenced documents?

Answer:

- Embedded → Store related data in the same document.
- Referenced → Store in separate collections using ObjectId.

💡 **Tip:** Embed for bounded small data, reference for unbounded/large data.

Q37. What are Capped Collections?

Answer: Fixed-size collections that overwrite oldest documents when full.

💡 **Tip:** Perfect for logs, caches, IoT time-series.

Q38. What is GridFS in MongoDB?

Answer: A specification for storing large files (>16 MB) by splitting into smaller chunks.

💡 **Tip:** Mention media storage (images, videos) as use case.

Q39. What is the 16MB document limit in MongoDB?

Answer: MongoDB restricts max document size to 16MB. For larger → use GridFS.

💡 **Tip:** Common trick question — always mention GridFS.

Q40. What is the difference between normalized and denormalized schema in MongoDB?

Answer:

- Normalized → references (less duplication, more joins).
- Denormalized → embedded/duplicated (fast reads, more storage).

💡 **Tip:** MongoDB often prefers denormalization for performance.

◆ Section 9: Replication & Sharding

Q41. What is a Replica Set in MongoDB?

Answer: Group of mongod instances with one primary (writes) and multiple secondaries (reads/failover).

💡 **Tip:** Interviewers love: “*Replication = high availability.*”

Q42. How does MongoDB handle failover?

Answer: If primary fails, secondaries hold an election to choose a new primary.

💡 **Tip:** Always mention automatic failover.

Q43. What is Sharding in MongoDB?

Answer: Horizontal scaling method → splits data across multiple shards using a shard key.

💡 **Tip:** Good for big data apps like e-commerce, logs.

Q44. How do you choose a Shard Key?

Answer: Pick a field with:

- High cardinality
- Even distribution
- Matches query patterns

💡 **Tip:** Bad shard key = hotspots + uneven load.

Q45. What is the difference between Replication and Sharding?

Answer:

- Replication → copies of same data across servers (HA).
- Sharding → distributes data across servers (scalability).

💡 **Tip:** Many candidates confuse both — highlight difference.

◆ Section 10: Advanced Indexing & Performance

Q46. What are Multikey Indexes?

Answer: Indexes built on array fields, creating an index entry for each

element.

💡 **Tip:** Useful for tags, categories.

Q47. What is a Text Index in MongoDB?

Answer: A **Text Index** is a special type of index in MongoDB that allows you to **search text inside string fields** (like full-text search).

```
db.articles.createIndex({ content: "text" })
```

```
db.articles.find({ $text: { $search: "MongoDB" } })
```

💡 **Tip:** Good for blogs, search bars.

Q48. What is a Geospatial Index?

Answer: Geospatial indexes (2d, 2dsphere) allow MongoDB to efficiently query **location-based data** (latitude/longitude).

Used for queries like *find nearby restaurants*.

Supports \$near, \$geoWithin.

💡 **Tip:** Mention Google Maps-like queries.

Q49. What is a Covered Query?

Answer: A query satisfied entirely by the index → no need to fetch documents.

💡 **Tip:** Saves I/O, improves performance.

Q50. How do you use explain() in MongoDB?

Answer: The .explain("executionStats") method shows **how MongoDB executes a query**.

Helps you check **index usage** and query efficiency.

```
db.users.find({ age: 25 }).explain("executionStats")
```

Shows query plan, index usage, scanned docs.

Output (important fields):

- **nReturned** → number of documents returned.
- **totalDocsExamined** → how many docs scanned.

- **indexUsed** → which index (if any) was used.

💡 **Tip:** Check nReturned vs totalDocsExamined.

◆ Section 11: Security & Transactions

Q51. How do you secure MongoDB in production?

Answer: Enable authentication, TLS, RBAC, encryption at rest, IP whitelisting.

💡 **Tip:** Mention “Never expose MongoDB to internet without auth.”

Q52. What is Role-Based Access Control (RBAC)?

Answer: MongoDB assigns roles (read, readWrite, dbAdmin) to users.

💡 **Tip:** Common enterprise feature.

Q53. What are Transactions in MongoDB?

Answer: Multi-document ACID transactions supported since v4.0.

💡 **Tip:** Mention important for financial apps.

Q54. What are Write Concerns in MongoDB?

Answer: Define acknowledgment level for writes (w:0,1,majority).

💡 **Tip:** Higher write concern = safer, slower.

Q55. What are Read Concerns in MongoDB?

Answer: Define consistency level for reads (local, majority, linearizable).

💡 **Tip:** Important in replicated environments.