

System Design Interview Q&As

◆ Section 1: Scalability & Basics

Q1. What is system design?

Answer: System design is about building scalable, reliable, and maintainable systems by balancing architecture, databases, and performance trade-offs.

Example: Designing Instagram with millions of users.

💡 Tip: Interviewers care about *your thought process*, not memorization.

Q2. What is scalability?

Answer: Scalability means handling growing load by adding more resources — vertical = bigger machine, horizontal = more machines.

Example: Amazon adding servers during Black Friday.

💡 Tip: Say *cloud-native apps prefer horizontal scaling*.

Q3. Difference between vertical vs horizontal scaling?

Answer: Vertical = upgrade CPU/RAM of one machine, Horizontal = add multiple servers. Horizontal is more fault-tolerant.

Example : Facebook uses horizontal scaling for billions of requests.

💡 Tip: Always mention horizontal scaling for modern apps.

Q4. What is high availability?

Answer: High availability ensures systems stay online with minimal downtime using replication and failover.

Example: Netflix is always available through multi-region failover.

💡 Tip: Say “HA = reliability + business continuity.”

Q5. What is fault tolerance?

Answer: Fault tolerance means the system continues working despite failures using redundancy and retries.

Example: Google Cloud runs across zones with auto-failover.

💡 Tip: “No single point of failure” is the key phrase.

◆ Section 2: Load Balancing

Q6. What is a load balancer?

Answer: A load balancer distributes traffic across servers to prevent overload and improve availability.

Example: AWS ELB routes traffic evenly.

💡 Tip: Mention performance + fault tolerance.

Q7. Difference between L4 vs L7 load balancers?

Answer: L4 balances traffic based on IP/TCP, while L7 works at HTTP/HTTPS and routes by content.

Example: L7 sends /images to image servers.

💡 Tip: Say “L7 = smarter, L4 = faster.”

Q8. What are sticky sessions in load balancers?

Answer: Sticky sessions tie users to the same server for session consistency.

Example: Shopping cart tied to one server.

💡 Tip: Mention drawback — “less scalable.”

Q9. How do DNS and load balancers work together?

Answer: DNS resolves domain → LB → routes traffic to servers.

Example: google.com → global DNS → LB → server.

💡 Tip: DNS + LB is the backbone of global services.

Q10. Global vs Local load balancing?

Answer: Local = distributes traffic in one data center, Global = across multiple regions.

Example: Cloudflare for global LB.

💡 Tip: “Global LB = geo-based routing.”

◆ Section 3: Caching

Q11. What is caching?

Answer: Caching stores frequently accessed data in fast storage to reduce response time.

Example: Redis cache for user sessions.

💡 Tip: Always link caching to *speed + cost savings*.

Q12. Client-side vs Server-side caching?

Answer: Client-side = browser/CDN, Server-side = DB query cache, Redis.

Example: Browser cache images; Redis caches queries.

💡 Tip: Show both for full answer.

Q13. What are cache eviction policies?

Answer: Methods to remove old data: LRU (Least Recently Used), LFU, FIFO.

Example: Redis uses LRU by default.

💡 Tip: Say “LRU = most common.”

Q14. CDN vs Cache?

Answer: CDN = caching at the edge (closer to users), cache = app/database level.

Example: Cloudflare CDN vs Redis cache.

💡 Tip: CDN = global reach, cache = local speed.

Q15. What is cache invalidation?

Answer: Removing stale data from cache using TTL, write-through, or write-back.

Example: Expiring old product prices.

💡 Tip: Mention “cache invalidation is hard but critical.”

◆ **Section 4: Messaging & Queues**

Q16. What is a message queue?

Answer: A queue decouples producers and consumers for async processing.

Example: Order service → Payment via RabbitMQ.

💡 Tip: Always connect queues with *scalability*.

Q17. Pub/Sub vs Message Queue?

Answer: Queue = one-to-one, Pub/Sub = one-to-many subscribers.

Example: Kafka for Pub/Sub in microservices.

💡 Tip: Say “Pub/Sub = broadcast, Queue = task handoff.”

Q18. What is backpressure in queues?

Answer: When consumers are slower than producers, messages pile up.

Example: Payment service lagging behind orders.

💡 Tip: Answer with “scale consumers or throttle producers.”

Q19. Dead letter queue (DLQ)?

Answer: DLQ stores failed/unprocessed messages for debugging.

Example: Invalid payment messages in DLQ.

💡 Tip: Say “DLQ improves reliability.”

Q20. Why use queues in system design?

Answer: Queues smooth spikes, allow retries, and enable async workflows.

Example: Email notifications via queue.

💡 Tip: “Queues = scalability + fault tolerance.”

◆ Section 5: Consistency & Reliability

Q21. What is CAP theorem?

Answer: In distributed systems, only 2 of Consistency, Availability, Partition Tolerance can be guaranteed.

Example: SQL = CA, DynamoDB = AP.

💡 Tip: Always mention trade-offs.

Q22. Strong vs Eventual consistency?

Answer: Strong = instant sync, Eventual = sync after some delay.

Example: SQL = strong, Amazon S3 = eventual.

💡 Tip: Use “shopping cart” as example.

Q23. What is replication?

Answer: Copying data across multiple servers for reliability and performance.

Example: Read replicas in MySQL.

💡 Tip: “Replication = high availability.”

Q24. What is sharding?

Answer: Splitting DB into smaller parts (shards) to scale horizontally.

Example: Users A–M in Shard1, N–Z in Shard2.

💡 Tip: “Sharding = scale but complex joins.”

Q25. What is quorum in distributed systems?

Answer: Minimum number of nodes needed to approve an operation.

Example: Raft/Paxos consensus.

💡 Tip: Say “quorum ensures consistency.”

◆ Section 6: Best Practices

Q26. How do you design for high traffic?

Answer: Use load balancers, caching, queues, and CDNs with horizontal scaling.

Example: Twitter during trending events.

💡 Tip: Always say “design for peak load.”

Q27. What is idempotency in APIs?

Answer: Same request multiple times = same result (no duplicate effect).

Example: Payment retry won't double charge.

💡 Tip: Mention *idempotency key*.

Q28. How do you monitor large-scale systems?

Answer: Use logging, metrics, and tracing tools like ELK, Prometheus, Grafana.

Example: Track API latency with Grafana.

💡 **Tip:** Always add *alerts + dashboards*.

Q29. What is blue-green deployment?

Answer: Two environments (Blue = live, Green = new). Switch traffic when new is stable.

Example: AWS CodeDeploy with Blue/Green.

💡 **Tip:** Say “safe CI/CD rollout.”

Q30. What are common bottlenecks in large systems?

Answer: DB queries, I/O latency, network bandwidth, poor architecture.

Example: Slow SQL query = app lag.

💡 **Tip:** Say “profiling + monitoring fix bottlenecks.”